



What are Sequence Learning Problem?

Introduction

- We have dealt with two types of networks:
 - feedforward neural networks and
 - convolution neural networks
-
- In both these networks the input was always of a fixed size.




- In feedforward and convolutional neural networks, **the size of the input was always fixed.**
- Further, **each input** to the network was **independent of the previous** or future inputs.
- For example, the computations, outputs and decisions for two successive images are completely independent of each other.





- In many applications the input is not of a fixed size.
- Further successive inputs may not be independent of each other.
- Consider the example of auto completion
- Given the first character 'd' you want to predict the next character 'e' and so on.

- Notice a few things
- First, **successive inputs are no longer independent** (while predicting 'e' you would want to know what the previous input was in addition to the current input).
- Second, **the length of the inputs and the number of predictions** you need to make is **not fixed** (for example, “learn”, “deep”, “machine” have different number of characters)
- Third, **each network** (orange-blue-green structure) **is performing the same task** (input: character , output : character)





- These are known as sequence learning problems.
- We need to look at a sequence of (dependent) inputs and produce an output (or outputs).
- Each input corresponds to one time step.






- Let us consider another example.
- Consider the task of predicting the part of speech tag (noun, adverb, adjective verb) of each word in a sentence.
- Once we see an adjective (social) we are almost sure that the next word should be a noun (man).
- Thus the current output (noun) depends on the current input as well as the previous input.
- Further the size of the input is not fixed (sentences could have arbitrary number of words)
- Notice that here we are interested in producing an output at each time step.
- Each network is performing the same task (input: word , output: tag).



- Sometimes we may not be interested in producing an output at every stage.
- Instead, we would look at the full sequence and then produce an output.
- For example, consider the task of predicting the polarity of a movie review.
- The prediction clearly doesn't depend only on the last word but also on some words which appear before.
- Here again we could think that the network is performing the same task at each step (input : word, output: +/-) but its just that we don't care about intermediate outputs.



- Sequences could be composed of anything (not just words)
- For example, a video could be treated as a sequence of images.

How to model these sequence learning problems?


- We look at something known as RNN : Recurrent Neural Networks
- Let's look at the function that is being executed at each time step:


$$s_i = \sigma(Ux_i + b)$$

$$y_i = \sigma(Vs_i + c)$$


$i = \text{timestep}$

- Since we want the same function to be executed at each timestep, we should share the same network (i.e., same parameters at each timestep)





- This parameter sharing also ensures that the network becomes agnostic to the length (size) of the input.
- Since we are simply going to compute the same function (with same parameters) at each timestep, the number of timesteps doesn't matter).
- We just create multiple copies of the network and execute them at each timestep.





- The function being computed at each time-step is as :



$$y_1 = f_1(x_1)$$

$$y_2 = f_2(x_1, x_2)$$

$$y_3 = f_3(x_1, x_2, x_3)$$



- The network is sensitive to the length of the sequence.
- For example, a sequence of length 10 will require 10 functions (f) whereas a sequence of length 100 will require 100 functions (f).



- The solution is to add a recurrent connection in the network.



$$s_i = \sigma(Ux + Ws_{i-1} + b)$$

$$y_i = \sigma(Vs_i + c)$$



or

$$y_i = f(x_i, s_i, W, U, V)$$

- S_i is the state of the network at timestep i .
- The parameters W, U, V, b which are shared across timesteps.
- The same network (and parameters) can be used to compute y_1, y_2, \dots, y_{100} .



.




Bidirectional Recurrent Neural Networks

- In sequence learning, so far we assumed that our goal is to model the next output given what we have seen so far, e.g., in the context of a time series or in the context of a language model.
- While this is a typical scenario, it is not the only one we might encounter.
- To illustrate the issue, consider the following three tasks of filling in the blank in a text sequence:
 - I am ____.
 - I am ____ hungry.
 - I am ____ hungry, and I can eat half a cake.

- I am ____.
- I am ____ hungry.
- I am ____ hungry, and I can eat half a cake.
- Depending on the amount of information available, we might fill in the blanks with very different words such as “happy”, “not”, and “very”.
- Clearly the end of the phrase (if available) conveys significant information about which word to pick.

- If we want to have a mechanism in **RNNs** that offers comparable **look-ahead ability**, we need to **modify the RNN design** that we have seen so far.
- Instead of **running an RNN only in the forward mode** starting from the first token, we start **another one from the last token running from back to front**.
- *Bidirectional RNNs* **add a hidden layer** that passes information **in a backward direction** to more flexibly process such information.




Architecture of a bidirectional RNN.

- For any time step t , given a minibatch input $\mathbf{X}_t \in \mathbb{R}^{n \times d}$ (number of examples: n , number of inputs in each example: d) and let the hidden layer activation function be ϕ .
- In the bidirectional architecture, we assume that the **forward and backward hidden states** for this time step are $\vec{\mathbf{H}}_t \in \mathbb{R}^{n \times h}$ and $\overleftarrow{\mathbf{H}}_t \in \mathbb{R}^{n \times h}$, respectively, where h is the number of hidden units.
- The **forward and backward hidden state updates** are as follows:

$$\begin{aligned}\vec{\mathbf{H}}_t &= \phi(\mathbf{X}_t \mathbf{W}_{xh}^{(f)} + \vec{\mathbf{H}}_{t-1} \mathbf{W}_{hh}^{(f)} + \mathbf{b}_h^{(f)}), \\ \overleftarrow{\mathbf{H}}_t &= \phi(\mathbf{X}_t \mathbf{W}_{xh}^{(b)} + \overleftarrow{\mathbf{H}}_{t+1} \mathbf{W}_{hh}^{(b)} + \mathbf{b}_h^{(b)}),\end{aligned}$$

- where the weights $\mathbf{W}_{xh}^{(f)} \in \mathbb{R}^{d \times h}$, $\mathbf{W}_{hh}^{(f)} \in \mathbb{R}^{h \times h}$, $\mathbf{W}_{xh}^{(b)} \in \mathbb{R}^{d \times h}$ and biases $\mathbf{b}_h^{(f)} \in \mathbb{R}^{1 \times h}$ and $\mathbf{b}_h^{(b)} \in \mathbb{R}^{1 \times h}$ are all the model parameters.




Architecture of a bidirectional RNN.

- Next, we concatenate the forward and backward hidden states $\vec{\mathbf{H}}_t$ and $\overleftarrow{\mathbf{H}}_t$ to obtain the hidden state $\mathbf{H}_t \in \mathbb{R}^{n \times 2h}$ to be fed into the output layer.
- In deep bidirectional RNNs with multiple hidden layers, such information is passed on as input to the next bidirectional layer.
- Last, the output layer computes the output (number of outputs: q):

$$\mathbf{O}_t = \mathbf{H}_t \mathbf{W}_{hq} + \mathbf{b}_q.$$

- Here, the weight matrix $\mathbf{W}_{hq} \in \mathbb{R}^{2h \times q}$ and the bias $\mathbf{b}_q \in \mathbb{R}^{1 \times q}$ are the model parameters of the output layer.



Architecture of a bidirectional RNN.

Deep RNN

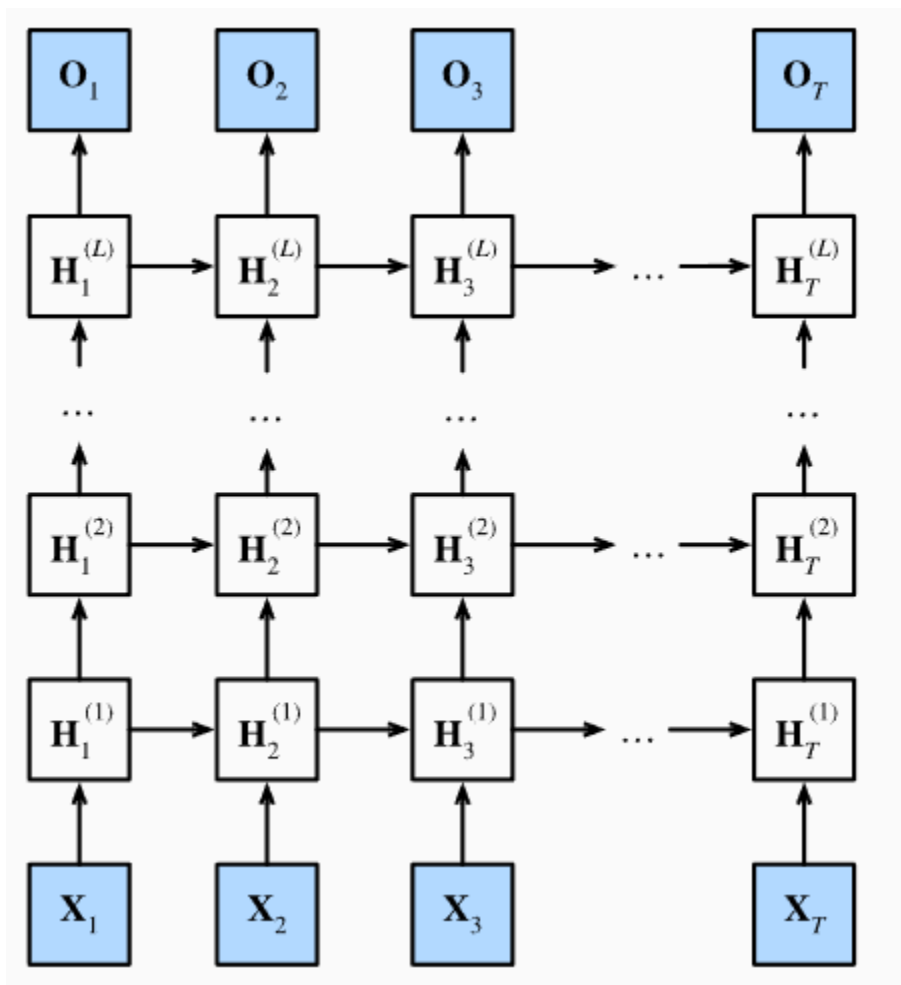


Fig : Architecture of a deep RNN.

Deep RNN Framework is a RNN framework for high-dimensional sequence problems like :

- Video classification
- Video future prediction

The figure shows a deep RNN with L hidden layers.

We could stack multiple layers of RNNs on top of each other.

This results in a flexible mechanism, due to the combination of several simple layers.

In particular, data might be relevant at different levels of the stack.

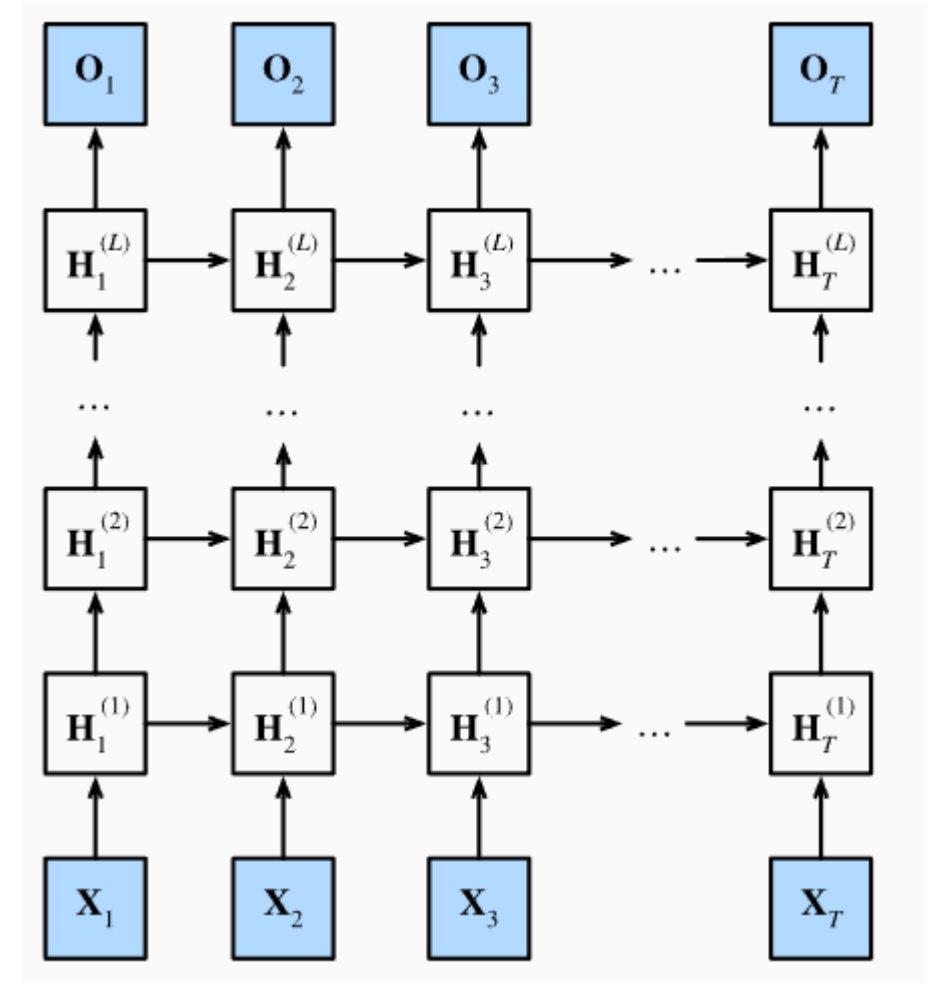
Each **hidden state is continuously passed to both the next time step of the current layer and the current time step of the next layer.**

Functional Dependencies

- Suppose that we have a minibatch input $\mathbf{X}_t \in \mathbb{R}^{n \times d}$ (number of examples: n , number of inputs in each example: d) at time step t .
- At the same time step, let the hidden state of the l^{th} hidden layer ($l=1, \dots, L$) be $\mathbf{H}_t^{(l)} \in \mathbb{R}^{n \times h}$ (number of hidden units: h) and the output layer variable be $\mathbf{O}_t \in \mathbb{R}^{n \times q}$ (number of outputs: q).
- Setting $\mathbf{H}_t^{(0)} = \mathbf{X}_t$, the hidden state of the l^{th} hidden layer that uses the activation function ϕ_l is expressed as follows:

$$\mathbf{H}_t^{(l)} = \phi_l(\mathbf{H}_t^{(l-1)} \mathbf{W}_{xh}^{(l)} + \mathbf{H}_{t-1}^{(l)} \mathbf{W}_{hh}^{(l)} + \mathbf{b}_h^{(l)}),$$

- where the weights $\mathbf{W}_{xh}^{(l)} \in \mathbb{R}_{h \times d}$ and $\mathbf{W}_{hh}^{(l)} \in \mathbb{R}_{h \times h}$, together with the bias $\mathbf{b}_h^{(l)} \in \mathbb{R}_{1 \times h}$, are the model parameters of the l^{th} hidden layer.



- In the end, the calculation of the output layer is only based on the hidden state of the final L^{th} hidden layer:

$$\mathbf{O}_t = \mathbf{H}_t^{(L)} \mathbf{W}_{hq} + \mathbf{b}_q$$

- where the weight $\mathbf{W}_{hq} \in \mathbb{R}_{h \times q}$ and the bias $\mathbf{b}_q \in \mathbb{R}_{1 \times q}$ are the model parameters of the output layer.
- Just as with MLPs, the number of hidden layers L and the number of hidden units h are hyperparameters. In other words, they can be tuned or specified by us.

